

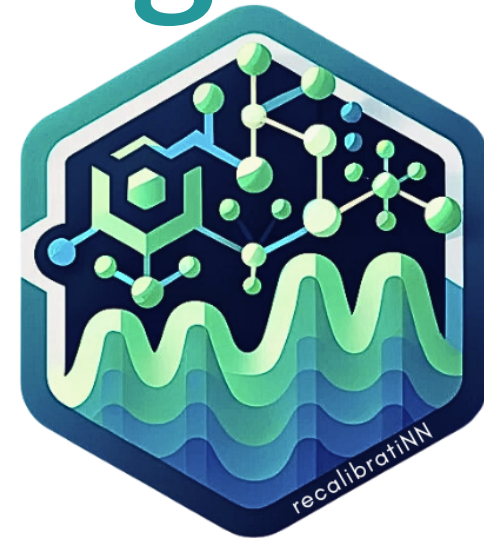
Recalibration of Gaussian Neural Networks regression models:

the *recalibratiNN* package

Carolina Musso

Instituto de Pesquisa e Estatística do
DF, Brazil

2024-07-11



A proper introduction

- Disclaimer: Me, the package and everything else.
 - Academic: Biological invasions, Fire Ecology, Ecotoxicology ...
 - Public Sector: Epidemiology, Sampling design and inference...
 - Free time: Bachelor degree in Statistics, Computational statistics, Bayesian methods, Neural Networks and Recalibration.
- R!
- Basically, I really wanted to develop a package.

Introduction: Neural Networks nowadays

- It should be able to quantify its uncertainty.
- NN can be constructed to produce probabilistic results:
 - Optimized by the log-likelihood.
 - Like any model, it can be miscalibrated.
 - A 95% CI should contain 95% of the true output.
 - $\mathbb{P}(Y \leq \hat{F}_Y^{-1}(p)) = p, \forall p \in [0, 1]$

Note

If optimized by MSE, I will be assuming a normal distribution.

Observing miscalibration

Consider a synthetic data set (x_i, y_i) , $i \in (1, \dots, n)$ generated by an heteroscedastic non-linear model:

$$x_i \sim \text{Uniform}(1, 10)$$

$$y_i | x_i \sim \text{Normal}(\mu = f_1(x_i), \sigma = f_2(x_i))$$

$$f_1(x) = 5x^2 + 10 ; f_2(x) = 30x$$

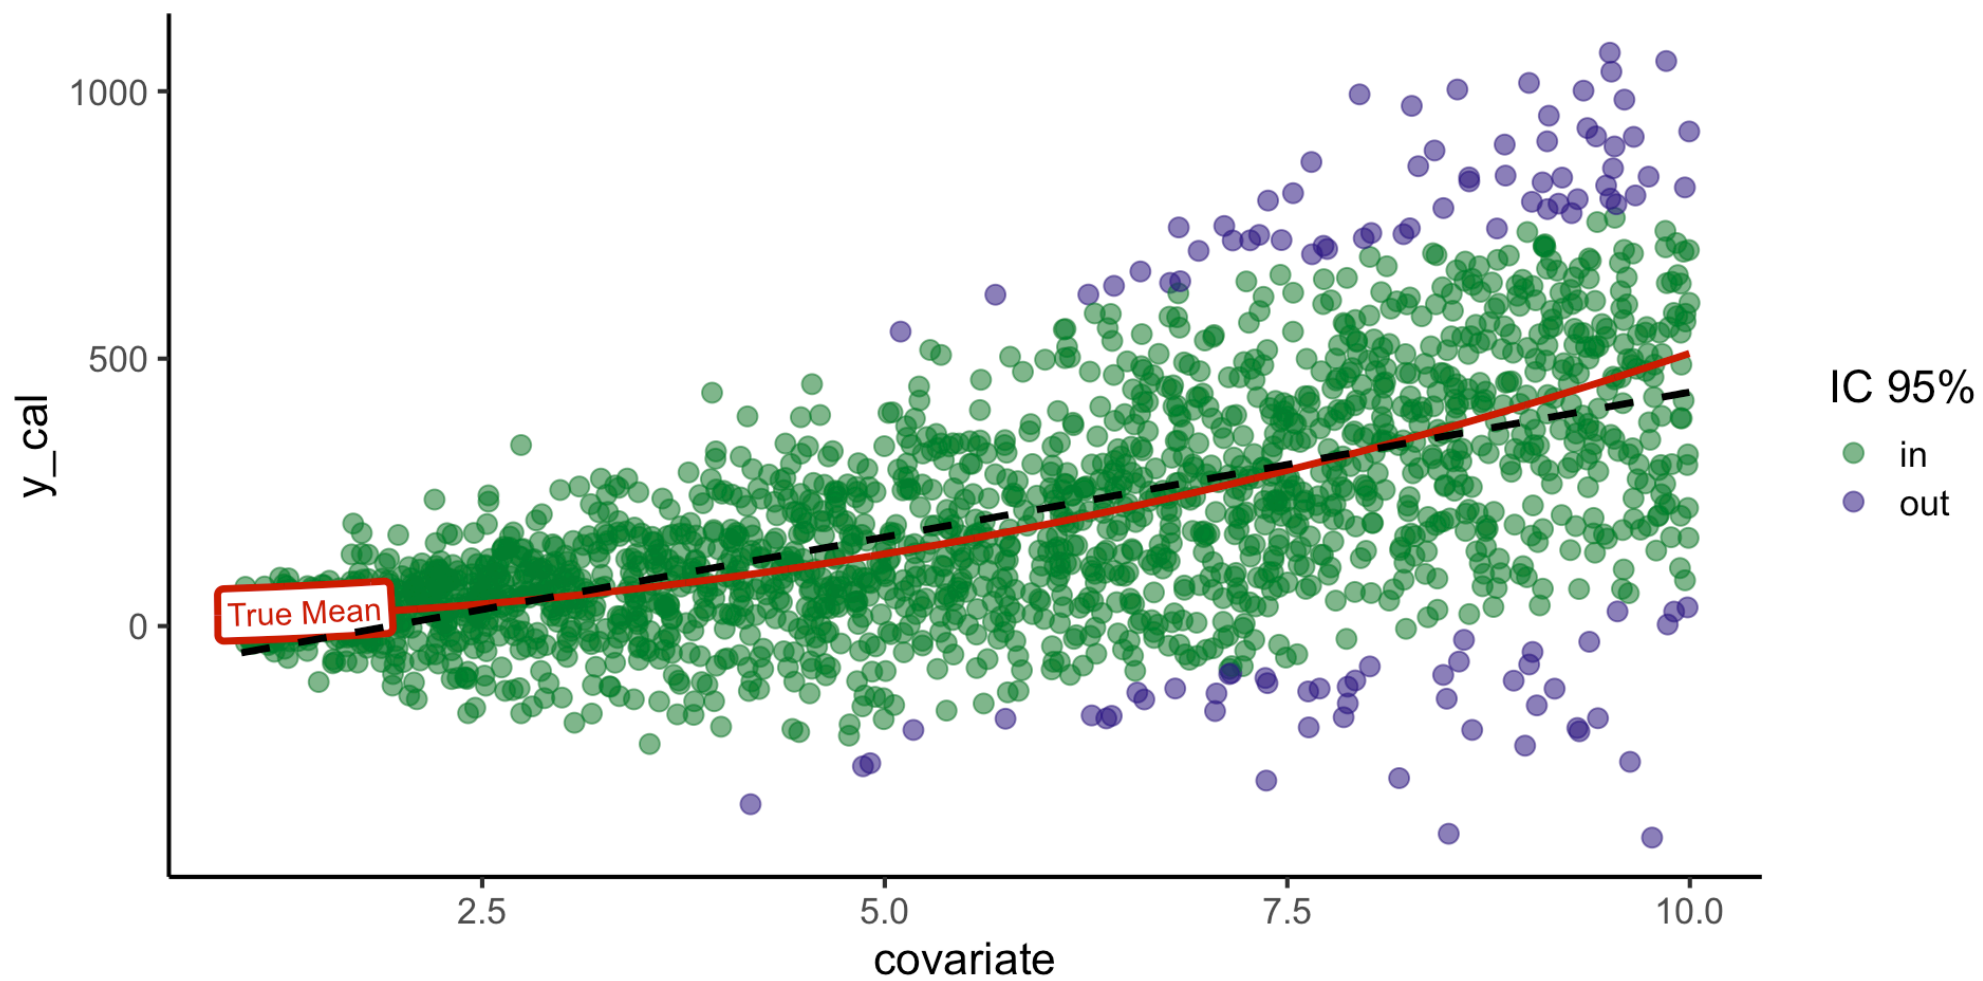
And the fitted model,

$$\hat{y}_i = \beta_0 + \beta_1 x_i + \epsilon_i, \epsilon_i \text{ iid } \sim N(0, \sigma)$$

Observing miscalibration

A simple linear regression, just to warm up.

- Global Coverage: 94.45%.



{fig-align='center' width=960 110%}

PIT - Values

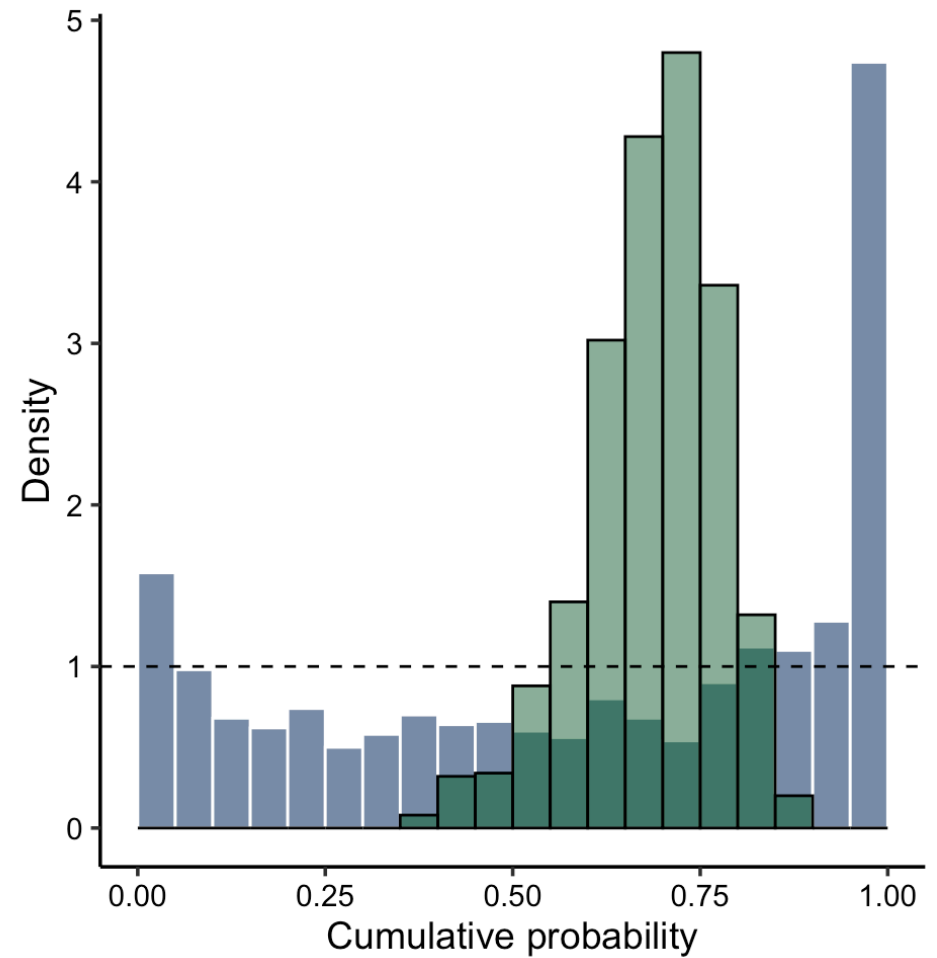
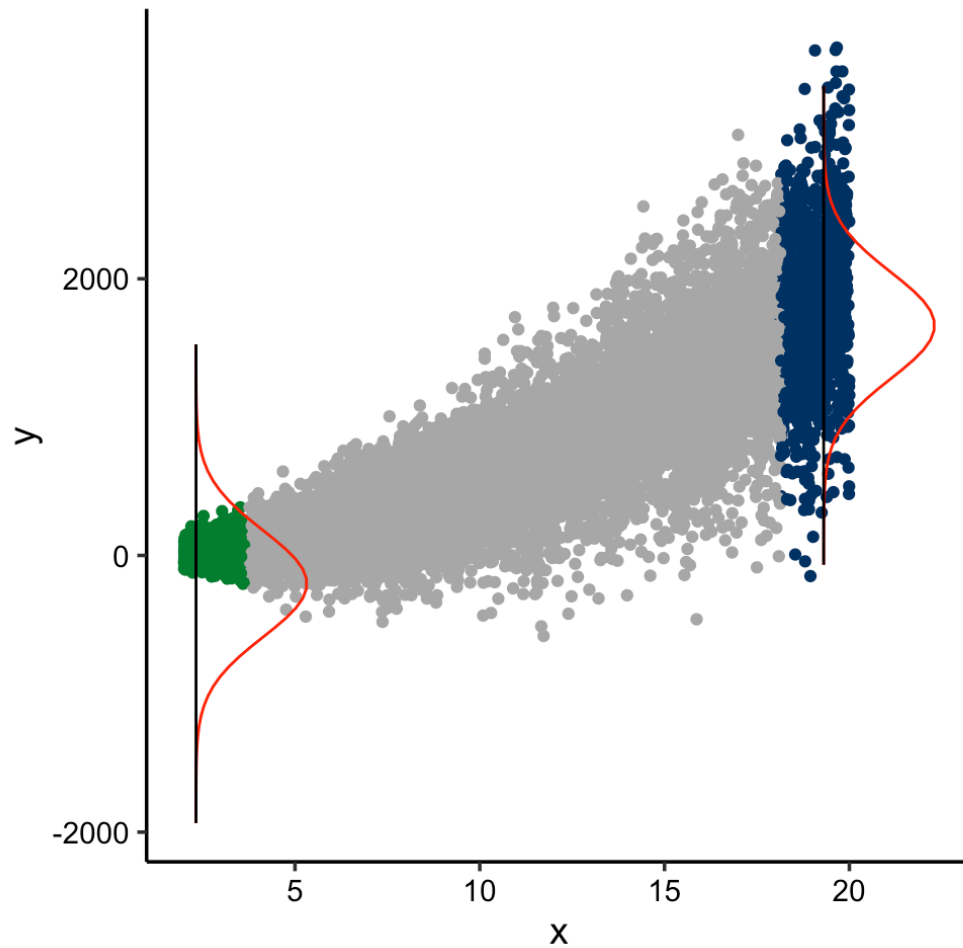
- Histogram of Probability Integral Transform (PIT) values.
- Let $F_Y(y)$ be the CDF of a continuous random variable Y , then:

$$U = F_Y(Y) \sim \text{Uniform}(0, 1)$$

- In particular, if $Y \sim \text{Normal}(\mu, \sigma)$:

$$Y = F_Y^{-1}(U) \sim \text{Normal}(\mu, \sigma)$$

Visualizing PIT-values



Recalibration

Available Packages

- R: [probably](#)
- Python: [ml_insights](#)
- Only global, focused on classification problems, and only applicable in covariate space.

Method:

- Torres et al (2024): Calibration across various representations of the covariate space: useful for Artificial Neural Networks (ANNs).

Algorithm

Algorithm 1 Torres *et al.* (2023) method implemented in the package

Input:

1:

- Recalibration set, $\{y_{\text{rec}}^{(i)}, \mathbf{x}_{\text{rec}}^{(i)}\}_{i=1}^n$, and new set, $\{\mathbf{x}_{\text{new}}^{(j)}\}_{j=1}^m$.
- A neural network and its associated predictive distribution, $\hat{F}(\cdot | \mathbf{X})$.
- A positive integer l defining the network's layer where the samples are to be compared.
- Neural network's outputs of the l -th layer on the recalibration set, $\{\mathbf{h}_{\text{rec}}^{(i)}\}_{i=1}^n$.
- A smoothing kernel $K_u(d)$ with scale parameter $u > 0$, which may be defined indirectly from a positive integer k that represents the number of observations to be used for recalibration.

Cumulative probabilities (PIT-values)

2: **for** $i \leftarrow 1$ **to** n **do**

3: Set $p_{\text{rec}}^{(i)} = \hat{F}_i(y_{\text{rec}}^{(i)} | \mathbf{x}_{\text{rec}}^{(i)})$.

4: **end for**

Recalibration

5: **for** $j \leftarrow 1$ **to** m **do**

6: Compute $\mathbf{h}_{\text{new}}^{(j)} = g(\mathbf{x}_{\text{new}}^{(j)})$, where g denotes the network's mapping to the l -th layer.

7: Apply the approximate KNN search method to identify the set of indices, I_j , corresponding to the observations in $\{y_{\text{rec}}^{(i)}, \mathbf{x}_{\text{rec}}^{(i)}\}_{i=1}^n$ for which $\|\mathbf{h}_{\text{rec}}^{(i)} - \mathbf{h}_{\text{new}}^{(j)}\|$ are within the k -smallest values.

8: **for** $i \in I_j$ **do**

9: Set $\tilde{y}_i^{(j)} = \hat{F}_j^{-1}(p_{\text{rec}}^{(i)} | \mathbf{x}_{\text{new}}^{(j)})$ and assign it a weight $w_i^{(j)} \propto K_u(\|\mathbf{h}_{\text{rec}}^{(i)} - \mathbf{h}_{\text{new}}^{(j)}\|)$.

10: **end for**

11: **end for**

Output:

12: A set of k weighted samples $\{(\tilde{y}_i^{(j)}, w_i^{(j)})\}$ from the recalibrated predictive distribution

$$\tilde{F}_j(\cdot | \mathbf{x}_{\text{new}}^{(j)}),$$

for $j = 1, \dots, m$.

The Package

- On GitHub
- and on CRAN

recalibratiNN package

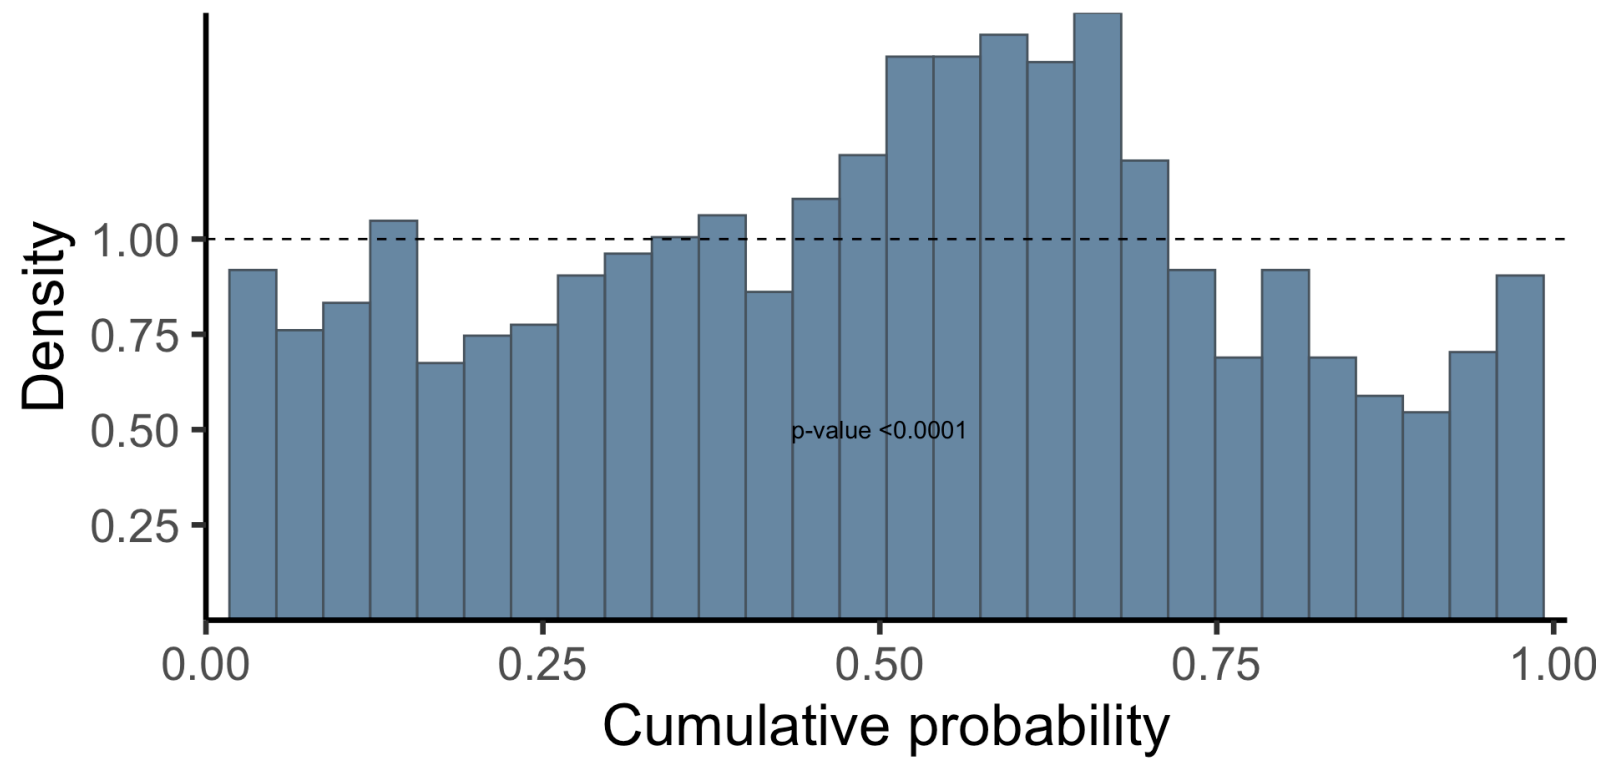
- 7 functions & 10 dependencies

Function	Description	Arguments
PIT_global	Calculates PIT values for the entire dataset	yca, yhat, mse
PIT_local	Calculates PIT values for each cluster	xca, yca, yhat, mse, clusters, p_neighbours, PIT
gg_PIT_global	Plots PIT values histogram	pit, type, fill, alpha, print_p
gg_PIT_local	Plots PIT values densities for kmeans clusters	pit_local, alpha, linewidth, pal, facet
recalibrate	Recalibrates the model	yhat_new, space_new, space_cal, pit_values, mse, type, p_neighbours, epsilon

Visualizing miscalibration


Global Calibration

```
1 pit <- PIT_global(y_cal = y_cal, # true values from calib. set.  
2                   y_hat = y_hat_cal, # predictions for calb. set.  
3                   mse    = MSE_cal) # MSE from calibration set.  
4  
5 gg_PIT_global(pit,  
6               type = "histogram",  
7               fill = "steelblue4",  
8               alpha = 0.8,  
9               print_p = TRUE  
10              )
```



Local Calibration

```
1 pit_local <- PIT_local(xcal = x_cal,  
2                       ycal = y_cal,  
3                       yhat = y_hat_cal,  
4                       mse = MSE_cal,  
5                       clusters = 6,  
6                       p_neighbours = 0.2,  
7                       PIT = PIT_global)  
8  
9 gg_PIT_local(pit_local)
```

Neural Networks

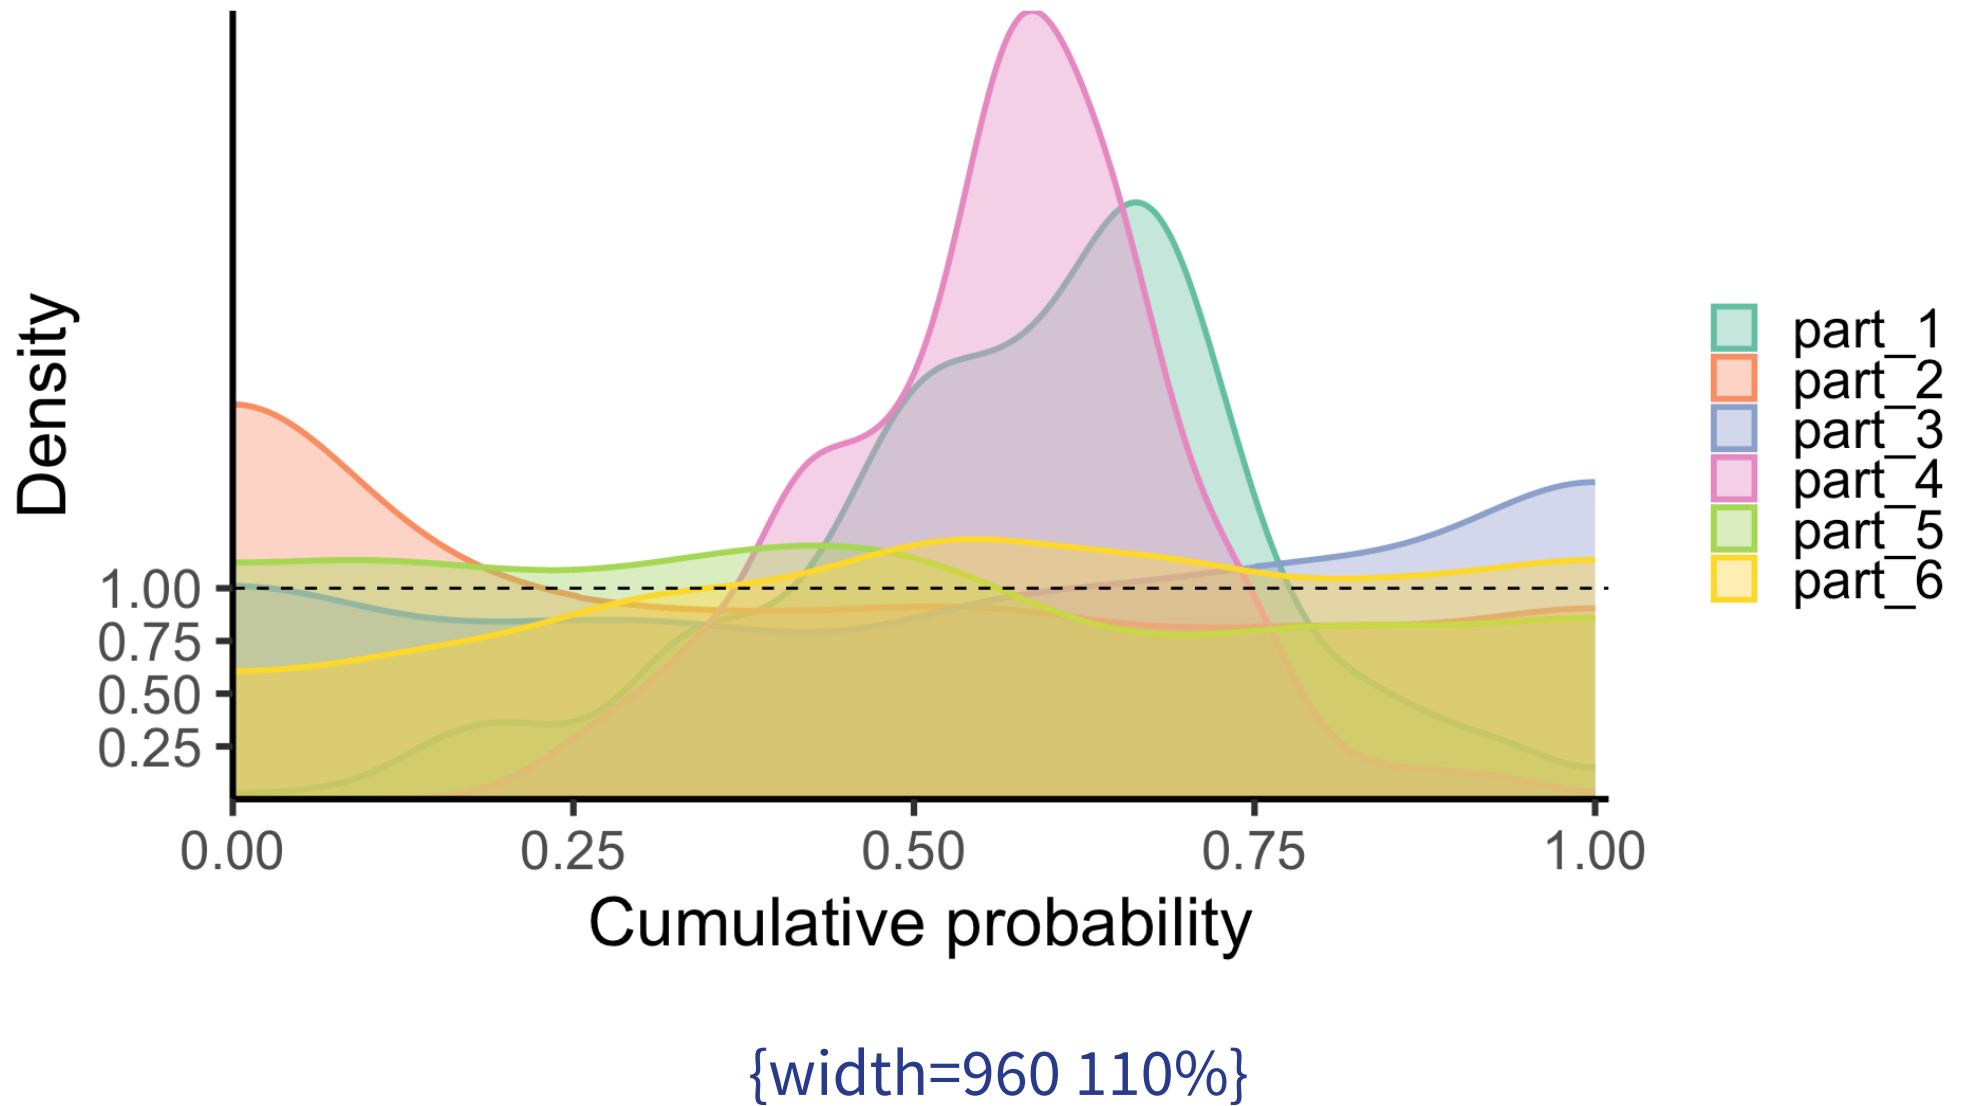
Data

```
1 set.seed(42)    # The Answer to the Ultimate Question of Life, The Universe, and
2
3 n <- 10000
4
5 x <- cbind(x1 = runif(n, -3, 3),
6            x2 = runif(n, -5, 5))
7
8 mu_fun <- function(x) {
9   abs(x[,1]^3 - 50*sin(x[,2]) + 30)}
10
11 mu <- mu_fun(x)
12 y <- rnorm(n,
13           mean = mu,
14           sd=20*(abs(x[,2]/(x[,1]+ 10))))
15
```

Keras

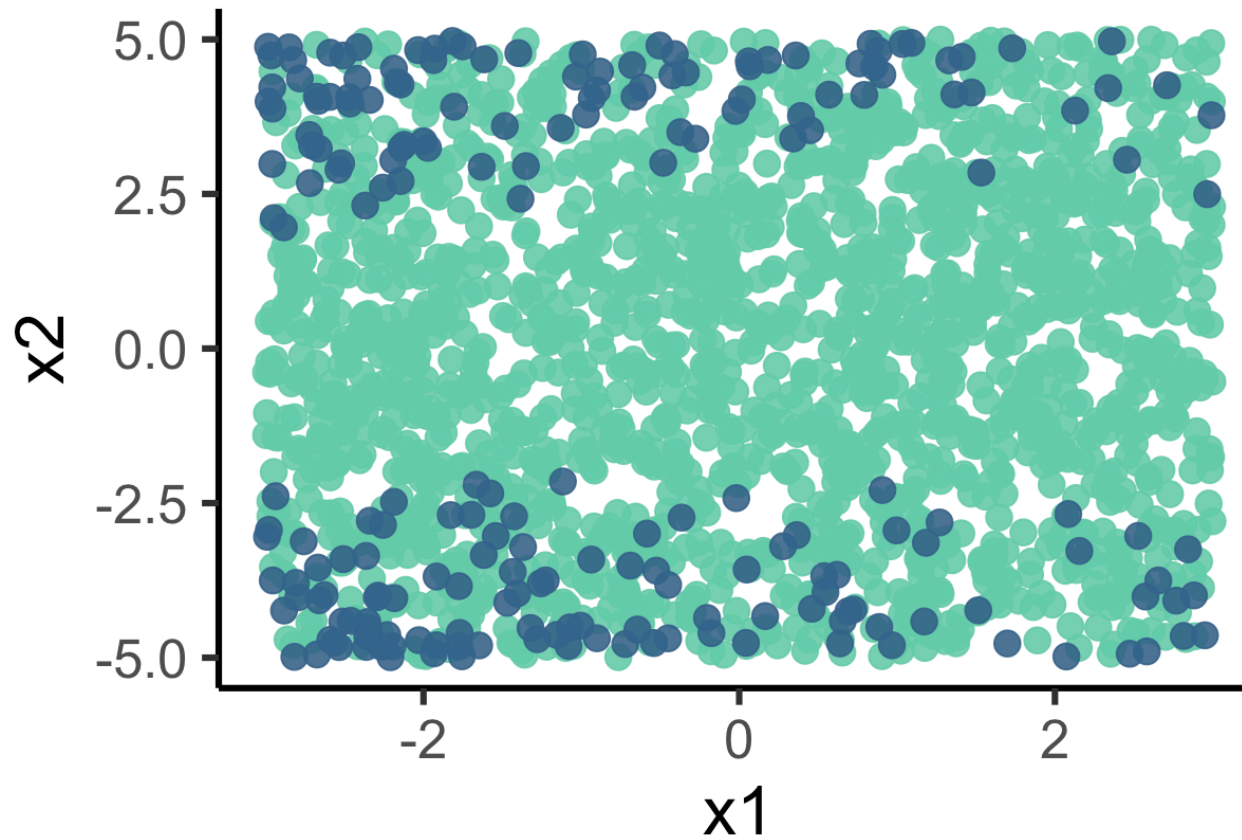
```
1 model_nn <- keras_model_sequential()  
2  
3 model_nn |>  
4   layer_dense(input_shape=2,  
5               units=800,  
6               use_bias=T,  
7               activation = "relu",  
8               kernel_initializer="random_normal",  
9               bias_initializer = "zeros") %>%  
10  layer_dropout(rate = 0.1) %>%  
11  layer_dense(units=800,  
12             use_bias=T,  
13             activation = "relu",  
14             kernel_initializer="random_normal",  
15             bias_initializer = "zeros") |>
```

Observing miscalibration



Coverage

Original coverage: 89.5 %



Confidence Interval

- in
- out

{width=960 110%}

Recalibration

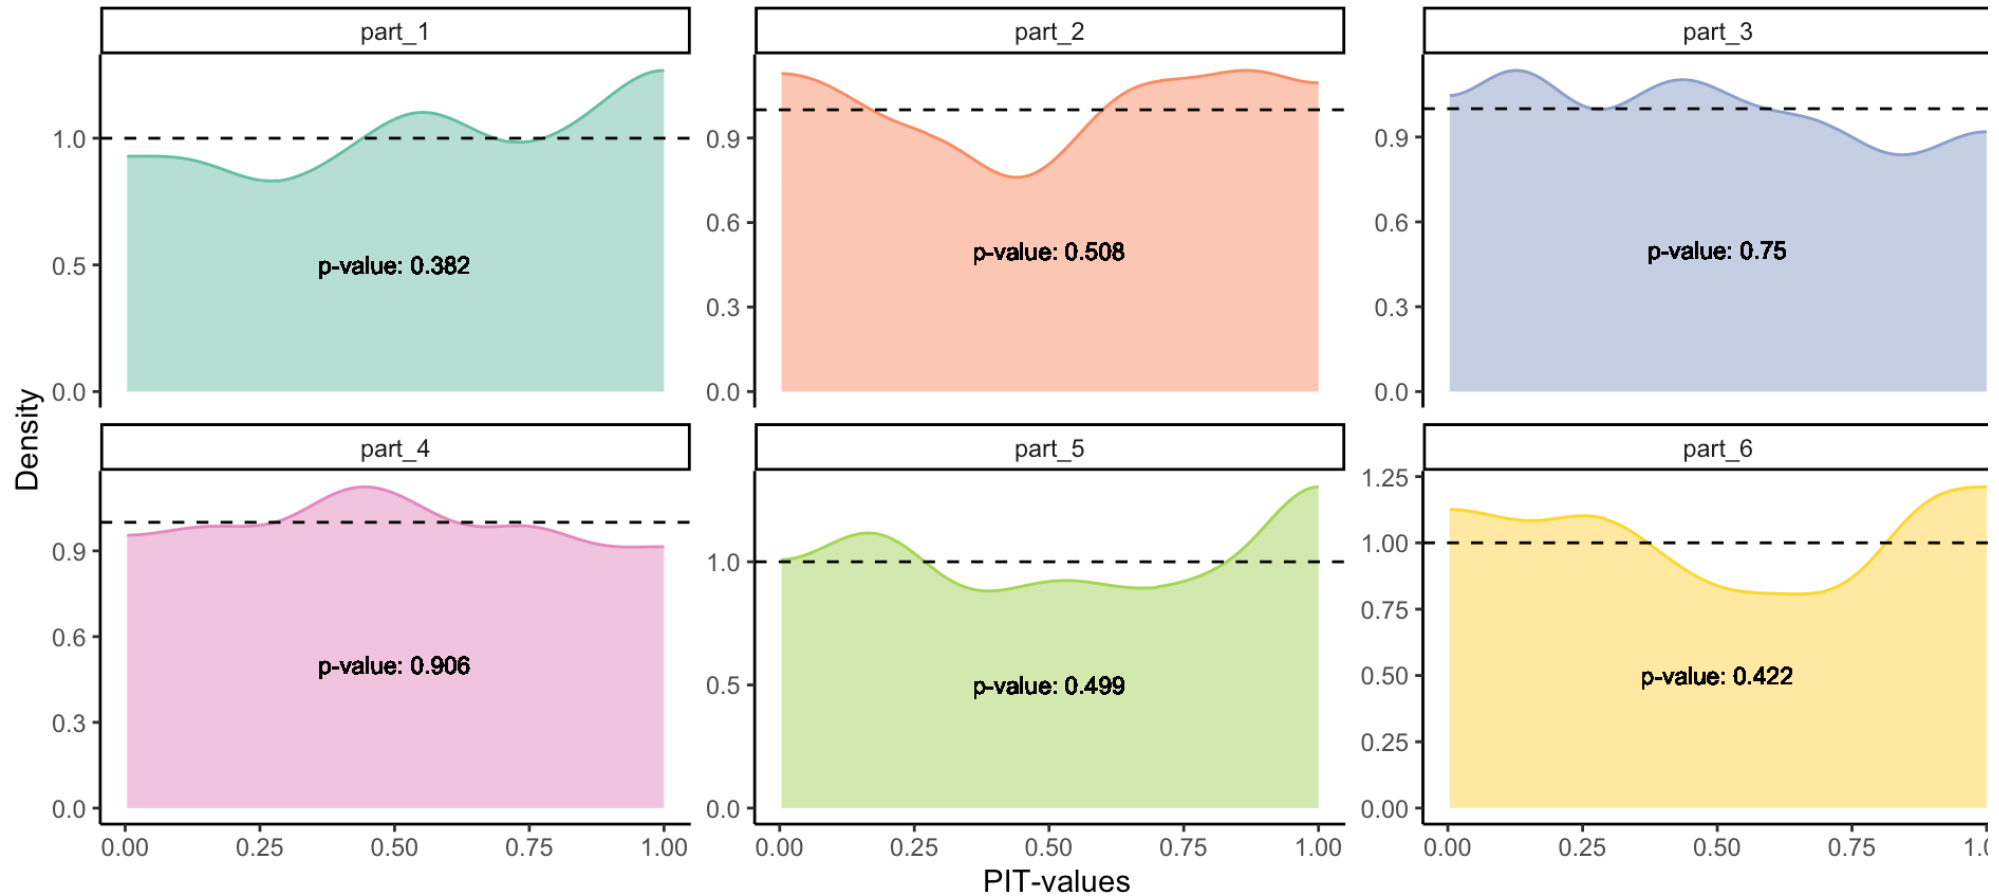
```
1 recalibrated <-  
2   recalibrate(  
3     pit_values = pit,      # global pit values calculated earlier.  
4     mse = MSE_cal,        # MSE from calibration set  
5     yhat_new = y_hat_test, # predictions of test set  
6     space_cal = x_cal,    # covariates of calibration set  
7     space_new = x_test,   # covariates of test set  
8  
9  
10  
11     type = "local",      # type of calibration  
12     p_neighbours = 0.08) # proportion of calibration to use as nearest neighbors  
13  
14 y_hat_rec <- recalibrated$y_samples_calibrated_wt
```

- That is it!
- These new values in `y_hat_rec` are, by definition, more calibrated than the original ones.

Shall we see?

After Local Calibration

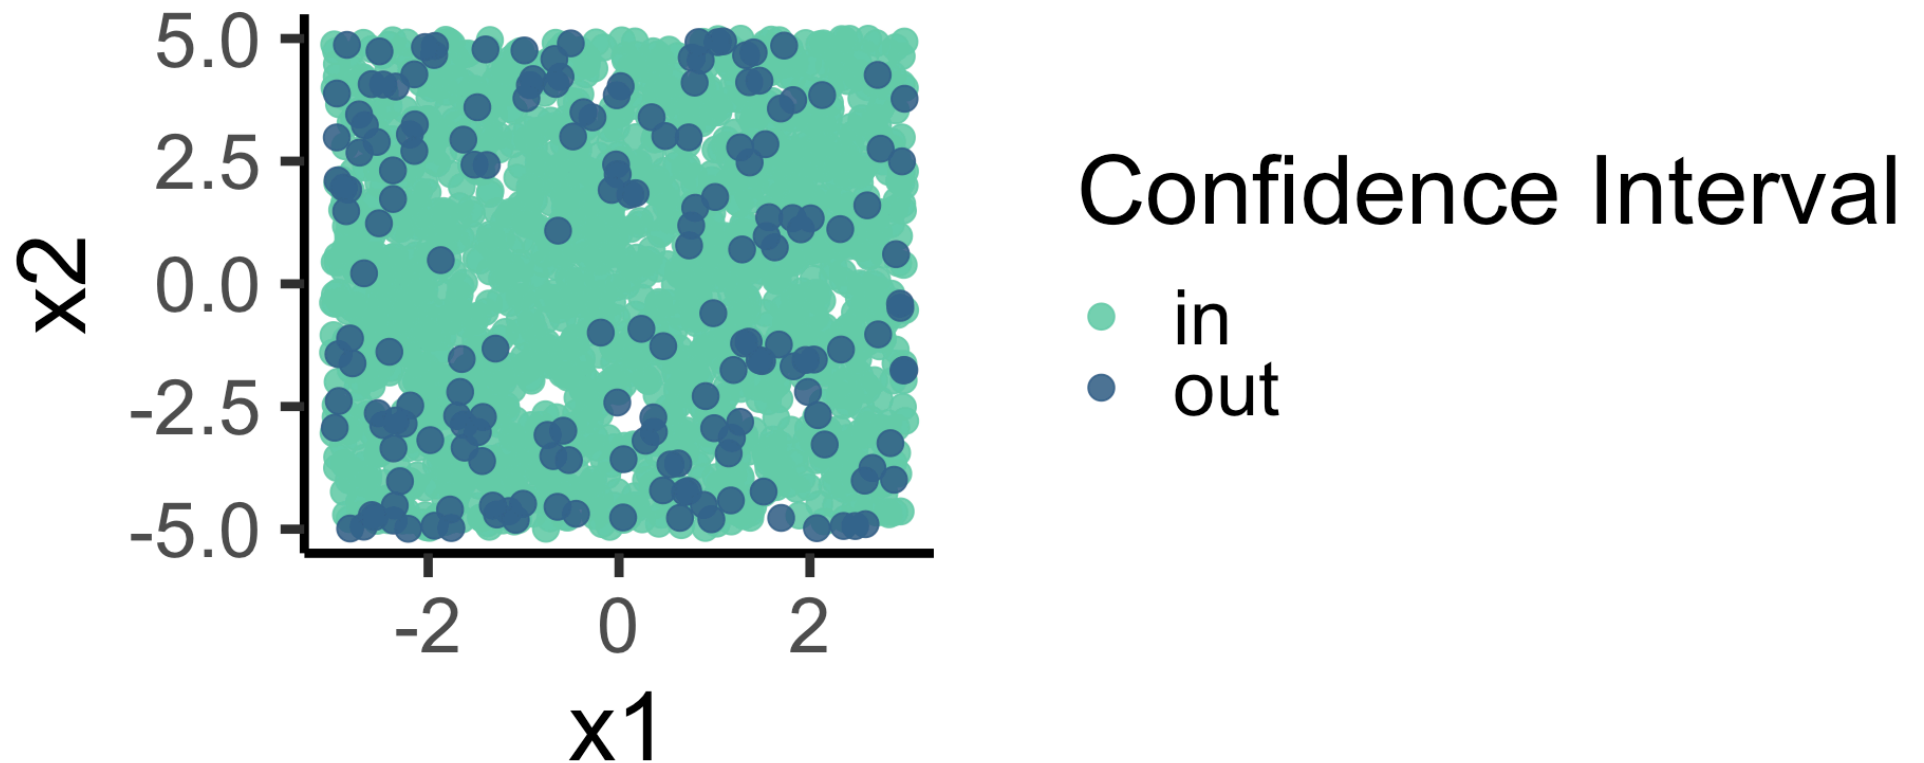
It looks so much better!!



{width=960 110}

Coverage

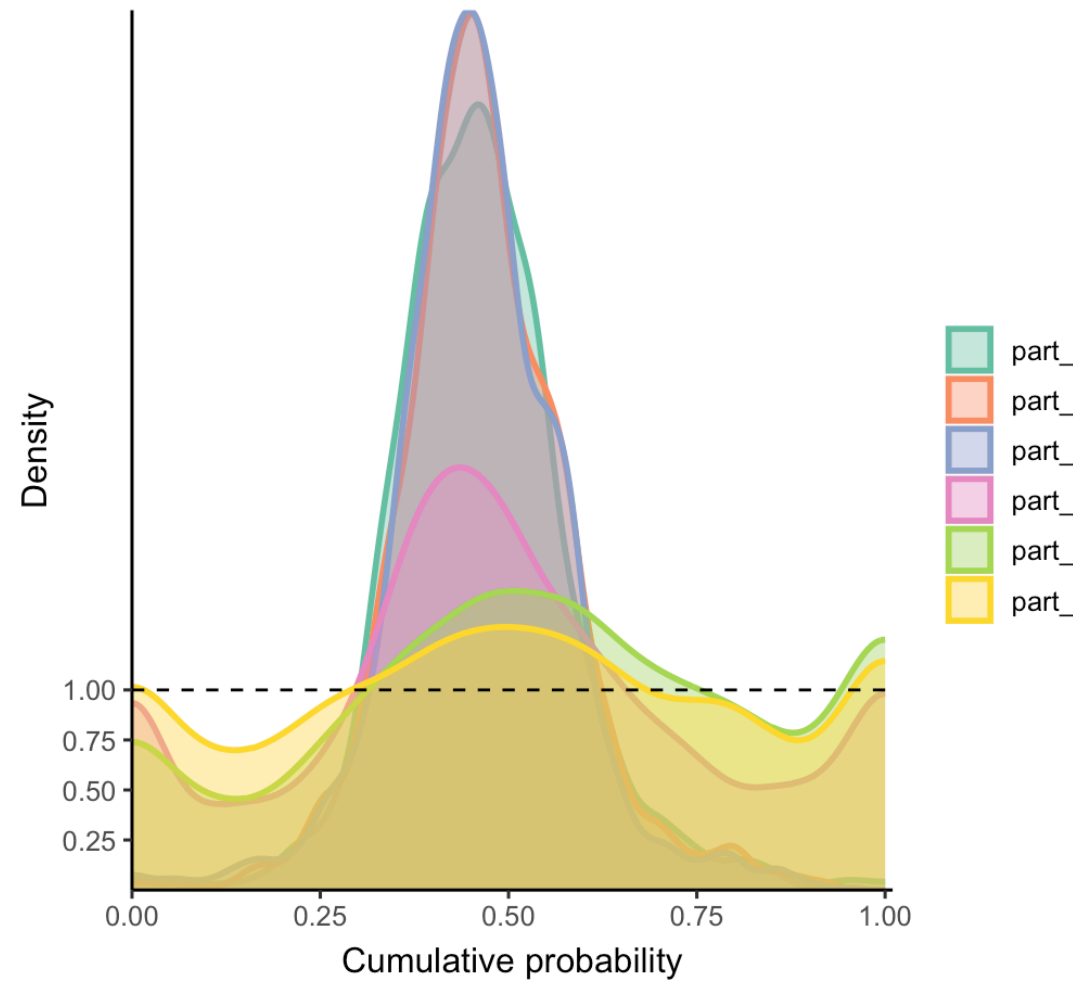
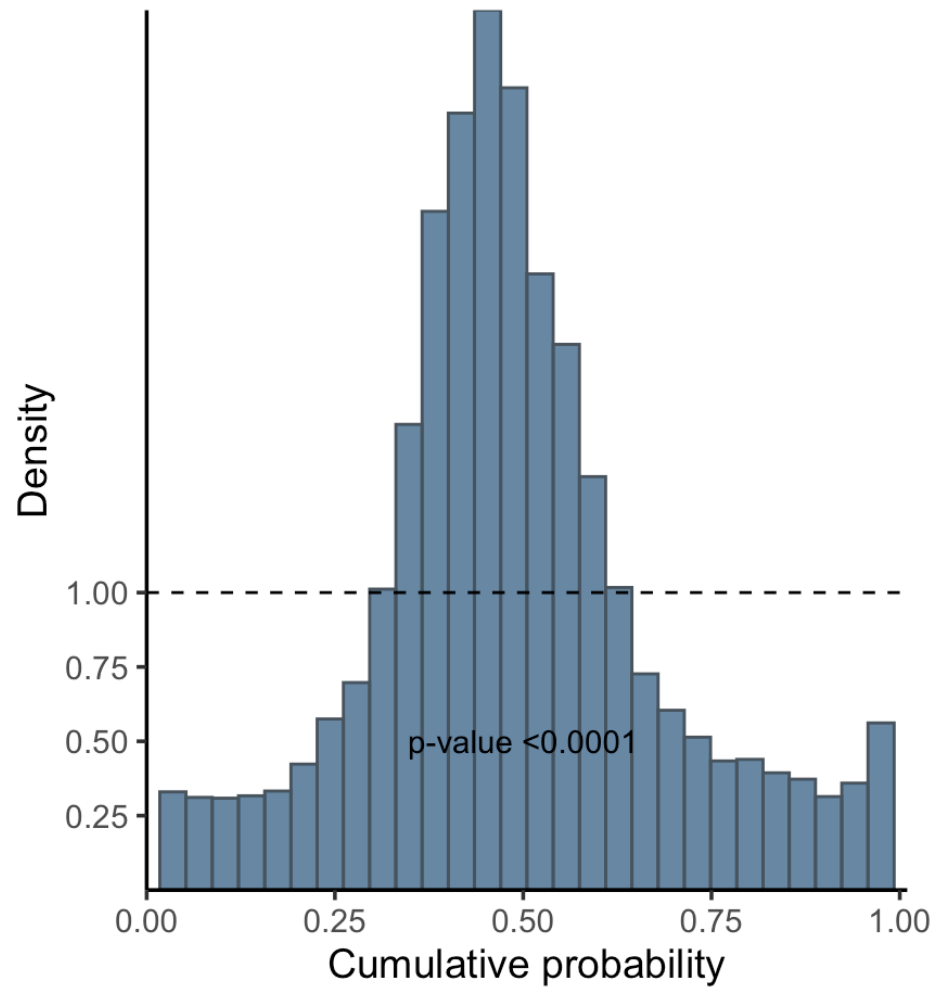
Recalibrated coverage: 90.8 %



{width=960 110}

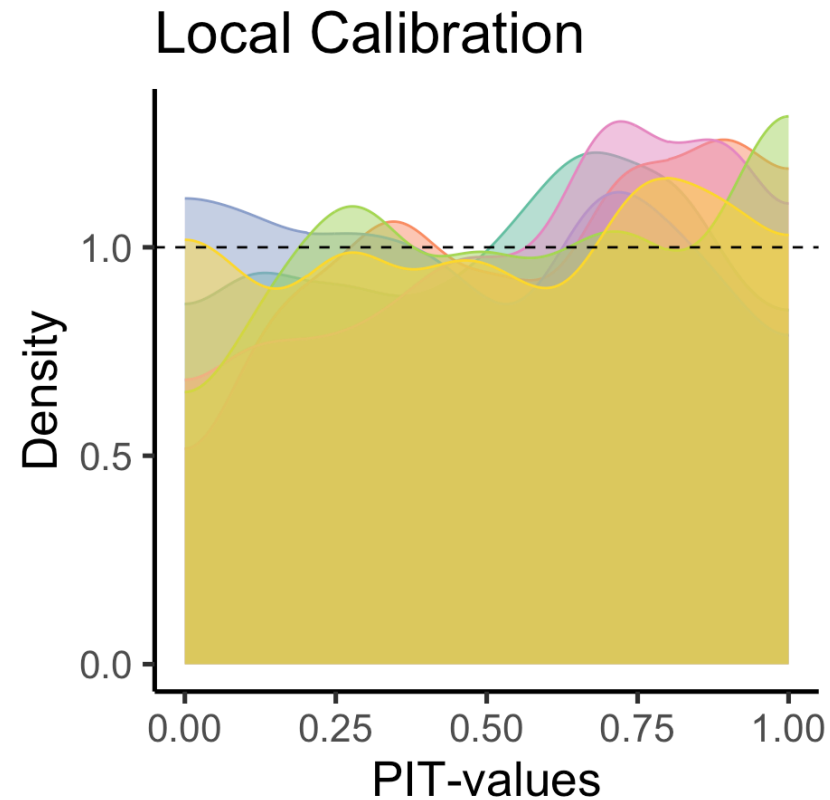
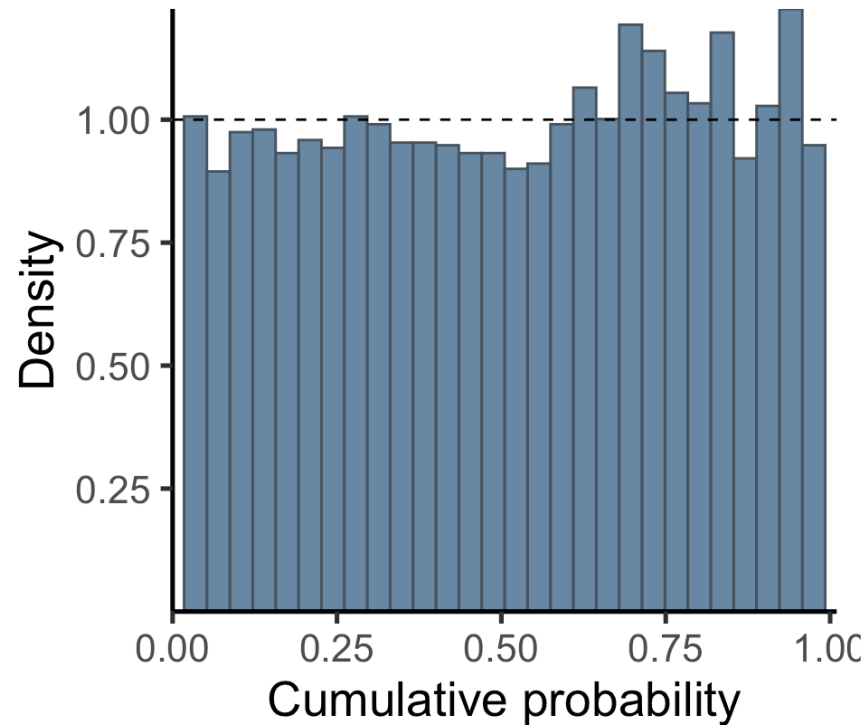
Real data

Diamonds dataset



After Recalibration

Calibrated using a second hidden layer.



Conclusions and Future Work

- Effective Visualization of Miscalibration.
- Advantages related to other packages
 - Focused in regression models
 - Local recalibration
 - Recalibration at intermediate layers.

Future Developments:

- Integration with other packages, broader input types, cross-validation methods
- Handle models with arbitrary predictive distributions.

Thank You!

[GitHub](